

CLOUD SECURITY

Zero to Hero

A practical course on securing AWS, Azure, and Google Cloud — from first principles to real attacks, container and DevSecOps defence, governance, and a graded interview question bank.



The path this course follows: each rung assumes the one below it.

Black & white edition · 60+ interview Q&A · vendor-neutral · nexroa

How to Use This Course

This course answers one question before all others: what actually changes about security when your servers, data, and networks move into someone else's data centre? The short answer is that you stop defending a perimeter you own and start defending an identity-and-configuration surface you share with a provider. Almost every cloud attack and defence in this course follows from that single shift.

The chapters are a ladder. Do not skip — each concept supports the next.

- Part I — Zero: what the cloud is, and the Shared Responsibility Model that governs everything.
- Part II — Novice: the three core domains — identity, network, and data.
- Part III — Practitioner: how real cloud breaches happen, and how to stop them.
- Part IV — Defender: containers, Kubernetes, serverless, and DevSecOps at scale.
- Part V — Hero: compliance, incident response, and running a cloud security program.
- Part VI — Interview Bank: 60+ graded questions with model answers.

READ FIRST: Vendor-neutral by design

AWS, Azure, and Google Cloud use different names for the same ideas (an AWS 'IAM Role' is close to an Azure 'Managed Identity' is close to a GCP 'Service Account'). This course teaches the concepts, then notes vendor terms where useful. Learn the concept once; the console names are just labels.

Contents

PART I — ZERO: FOUNDATIONS

1. What the Cloud Actually Is
2. The Shared Responsibility Model
3. The Cloud Threat Landscape

PART II — NOVICE: THE CORE DOMAINS

4. Identity & Access Management (IAM)
5. Network Security in the Cloud
6. Data Security & Encryption

PART III — PRACTITIONER: ATTACKS & DEFENCE

7. How Cloud Breaches Actually Happen
8. Misconfiguration: the Number-One Cause
9. Secrets Management
10. Logging, Monitoring & Detection

PART IV — DEFENDER: SCALE & AUTOMATION

11. Container & Kubernetes Security
12. Serverless Security
13. DevSecOps & Infrastructure as Code

PART V — HERO: GOVERN & OPERATE

14. Compliance & Frameworks
15. Cloud Incident Response
16. Building a Cloud Security Program

PART VI — INTERVIEW PREPARATION

17. Interview Q&A Bank (60+ graded)
18. Course vs. Q&A: Which Matters More?

Appendix A — 30-Day Study Plan

Appendix B — Glossary

Page numbers omitted here so the contents stay accurate as the course is revised; every chapter starts on its own labelled page and appears in running headers.

PART I**Zero: Foundations**

“You cannot secure what you do not understand. First, understand what changed when the servers stopped being yours.”

1. What the Cloud Actually Is

Strip away the marketing and ‘the cloud’ is one thing: renting someone else’s computers, on demand, over the internet. Instead of buying servers and locking them in a room, you rent compute, storage, and networking from a provider (AWS, Microsoft Azure, Google Cloud) and pay for what you use. Everything about cloud security follows from that ownership change.

The three service models

How much you rent determines how much you must secure. This is the single most important distinction in the field.

Model	You rent...	Example	You still secure
IaaS	Raw compute / storage / network	AWS EC2, Azure VMs	OS, apps, data, access
PaaS	A managed platform to run apps	App Engine, Azure App Service	Your code, data, access
SaaS	Finished software	Gmail, Salesforce	Your data & who can access it

Why cloud security is not just “IT security elsewhere”

- No physical perimeter: there is no office wall or firewall at the edge. Your resources are reachable over the internet by default unless you say otherwise.
- Everything is an API: every resource is created, changed, and deleted through an API call. That means one leaked key can rewrite your whole environment in seconds.
- Speed and scale: a single command can launch a thousand servers — or expose a thousand. Mistakes replicate as fast as features.
- Identity is the perimeter: since anyone on the internet can reach the API, who you are (identity) replaces where you are (network location) as the primary control.

MENTAL MODEL: First principle

On-premises, an attacker had to get through your network to reach a server. In the cloud, the server is already on the internet; the attacker just needs valid credentials or a misconfiguration. The battle moved from the network edge to identity and configuration.

2. The Shared Responsibility Model

This is the foundation the entire field rests on. Get it wrong and you will secure the wrong things. The model answers one question: in any cloud service, who is responsible for securing what — you or the provider?

The rule of thumb: the provider secures the cloud itself (their hardware, data centres, hypervisors); you secure what you put in the cloud (your data, access, configuration). But the exact line moves depending on the service model — and that is where people get breached.

The Shared Responsibility Model

	On-Premises	IaaS	PaaS	SaaS
Data & Access	C	C	C	C
Application	C	C	C	P
Runtime	C	C	P	P
Operating System	C	C	P	P
Virtualization	C	P	P	P
Servers / Compute	C	P	P	P
Storage	C	P	P	P
Networking	C	P	P	P
Physical Facility	C	P	P	P

C = Customer secures
 P = Provider secures

Who secures what, by service model. As you move IaaS → SaaS the provider takes on more, but you always own data and access. The customer's share never reaches zero.

The dangerous misreadings

- “It’s in the cloud, so it’s secure.” False. The provider secures their layer; your misconfigured bucket is entirely your problem.
- “SaaS means nothing to secure.” False. You still control who can access the data and how it is shared — the most common SaaS breach is a badly shared file or an account without MFA.
- “The provider backs up my data.” Usually false by default. Availability of the platform is theirs; your data’s backup and integrity is often yours.

RULE: The line you always own

Across every model — IaaS, PaaS, SaaS — two things are always your responsibility: your data and who can access it. If you remember nothing else about shared responsibility, remember that identity and data never become the provider’s job.

3. The Cloud Threat Landscape

A threat model answers four questions: what are we protecting, who attacks it, how, and what happens if they win. In the cloud, the surprising answers are that the attacker usually gets in with a valid credential (not an exploit), and the damage is usually data exfiltration at scale.

The Cloud Attack Kill Chain



Most cloud breaches begin at Initial Access via a leaked credential or a misconfiguration — not a zero-day.

The cloud kill chain. Note where it starts: most breaches begin with a leaked credential or a misconfiguration, not a sophisticated zero-day.

Who attacks cloud environments, and why

Actor	Typical goal	Typical method
Opportunistic scanners	Anything exposed	Scan the internet for public buckets, open ports, leaked keys
Financially motivated	Data to sell, ransom, or mine crypto	Phished/leaked credentials, then privilege escalation
Insiders	Data theft or sabotage	Abuse of legitimate, over-broad access
Advanced / targeted	Specific data or persistence	Chained misconfigs, supply chain, living off the cloud

What you are protecting (the CIA triad, cloud-flavoured)

- Confidentiality: customer data, secrets, source code, personal information (PII).
- Integrity: the correctness of your data and infrastructure — an attacker who alters config or records is as dangerous as one who steals.
- Availability: uptime and cost. In the cloud, availability includes financial availability — an attacker can run up a crippling bill (crypto-mining, denial-of-wallet).

PART II


Novice: The Core Domains

“Three domains carry most of cloud security: who can act (identity), what can talk to what (network), and what protects the data itself.”

4. Identity & Access Management (IAM)

If you master one domain, make it this one. In the cloud, identity is the perimeter. Most breaches are ultimately an identity failure: a leaked key, an over-privileged role, a missing MFA. IAM is the system that decides who (identity) can do what (permission) to which resource.

Identity is the New Perimeter



The IAM triangle. An identity assumes a policy that grants scoped access to a resource. Least privilege means the policy grants only the exact actions needed — nothing more.

The core building blocks

Concept	What it is	Vendor terms
Principal / Identity	A user, group, application, or workload that acts	IAM user, service account, managed identity
Role	A set of permissions an identity temporarily assumes	AWS role, Azure role, GCP role
Policy	The rules stating what actions are allowed on what	IAM policy, RBAC assignment
MFA	A second proof of identity beyond a password	MFA / 2FA

The four IAM rules that prevent most breaches

- Least privilege: grant the minimum permissions the task needs. Start from zero and add, never start from admin and trim.
- MFA everywhere, always on privileged accounts: a password alone is one leak away from disaster.
- Prefer roles over long-lived keys: temporary, auto-rotating credentials beat static access keys that live in a file forever.
- No standing admin: humans should not hold permanent god-mode. Grant elevated access just-in-time, for a task, then revoke.

ANTI-PATTERN: The most dangerous IAM word: “*”

A policy that grants Action: * on Resource: * means “this identity can do anything to everything.” Attach it to a workload with an internet-facing bug and you have handed the attacker your entire account. Wildcards are where least privilege goes to die — hunt them down.

5. Network Security in the Cloud

Even though identity is the new perimeter, the network still matters — as a second wall. Cloud network security is about controlling what can talk to what. The goal is that even if one component is compromised, the attacker cannot freely reach everything else (this is segmentation, and it limits lateral movement).

The building blocks

Concept	What it does	Vendor terms
Virtual network	Your private, isolated network in the cloud	VPC (AWS/GCP), VNet (Azure)
Subnet	A segment of that network; public or private	Subnet
Security group	A stateful firewall around a resource	Security group, NSG
Private endpoint	Reach a managed service without the public internet	PrivateLink, Private Endpoint

Principles that matter

- Default deny: allow only the traffic you explicitly need; block everything else. Never leave management ports (SSH 22, RDP 3389) open to the whole internet.
- Public vs private subnets: only put things that must be reachable (a load balancer) in public subnets; keep databases and internal services private.
- Segment to limit blast radius: separate environments (prod/dev) and tiers (web/app/data) so a breach in one cannot trivially reach the others.
- Prefer private connectivity: keep traffic to managed services (databases, storage) off the public internet using private endpoints.

MODERN MODEL: Zero Trust in one line

“Never trust, always verify.” Do not assume a request is safe just because it comes from inside your network. Verify identity and authorisation on every request, everywhere. In the cloud, where the network edge is fuzzy, this mindset is not optional.

6. Data Security & Encryption

Ultimately, attackers want the data. Every other control exists to protect it. Data security has two jobs: control access to the data (IAM, again) and protect the data itself so that even if someone gets the bytes, they are useless without the key.

Encryption: the two states

- Encryption at rest: data stored on disk is encrypted, so a stolen disk or storage object is unreadable. In the cloud this is often on by default — but you must verify it.
- Encryption in transit: data moving over the network is encrypted (TLS), so it cannot be read if intercepted. Enforce TLS everywhere; reject unencrypted connections.

Keys are the real secret

Encryption only moves the problem: now the key is what must be protected. If the attacker gets the key, encryption bought you nothing. Cloud key management services (KMS) store and control keys, log every use, and let you rotate and revoke them.

Approach	Who holds the key	Trade-off
Provider-managed keys	The cloud provider	Easiest; you trust the provider fully
Customer-managed keys (KMS)	You, in the provider's KMS	More control, rotation, audit; small effort
Customer-supplied keys	You, entirely outside	Maximum control; maximum operational burden

RULE: Design rule

Never rely on encryption alone. Encryption protects against a stolen disk; it does nothing against a valid credential that is simply allowed to read the data. Access control (IAM) and encryption are partners — you need both. Most data breaches are access failures, not broken crypto.

Preventing accidental exposure

- Block public access on storage by default; the classic breach is a bucket set to 'public'.
- Classify data (public / internal / confidential / regulated) so controls match sensitivity.
- Redact or tokenise sensitive fields (PII) where the full value is not needed.
- Turn on data-access logging so you can see who read what, when.

PART III

Practitioner: Attacks & Defence


“Theory earns respect; knowing how breaches actually unfold earns a job. This is where the course gets concrete.”

7. How Cloud Breaches Actually Happen

Forget Hollywood. Most cloud breaches are not clever zero-days; they are boring, preventable chains: a leaked key or misconfiguration grants access, an over-privileged role lets the attacker escalate, and weak segmentation lets them reach the data. Understanding one canonical chain teaches the whole pattern.

The canonical attack: SSRF to credential theft

This is the pattern behind several of the largest cloud breaches on record. Follow it step by step; every defence in this course maps back to breaking one of these links.



A server-side request forgery (SSRF) bug lets an attacker trick the application into fetching the cloud metadata endpoint, which hands out the workload's temporary IAM credentials. Those credentials then unlock whatever the over-privileged role could reach.

Why each link exists — and how to break it


Link in the chain	Root cause	The defence
1. SSRF in the app	Unvalidated user-supplied URL	Input validation; block internal IP ranges
2. Metadata reachable	Legacy metadata service, no session	Enforce hardened metadata (e.g. IMDSv2)
3. Creds are powerful	Over-privileged workload role	Least privilege on the role
4. Data reachable	Flat network, broad access	Segmentation; scoped resource policies
5. Silent exfiltration	No monitoring on data access	Logging + anomaly detection + alerts

KEY IDEA: The lesson of the chain

No single control would have stopped every breach — but any one of these five defences would have broken the chain. That is defence in depth in action: you assume one control fails and make sure the next one holds. Least privilege is the highest-leverage link, because it caps what a stolen credential can do.

8. Misconfiguration: the Number-One Cause

If you remember one statistic: the large majority of cloud data breaches are caused by customer misconfiguration, not provider failure. The cloud gives you thousands of settings; a single wrong one can expose everything. This is the direct, day-to-day work of cloud security.



The recurring culprits. These are not exotic — they are default-adjacent mistakes that scanners find within minutes of exposure. Bar length is an illustrative index, not a formal ranking.

The classics, and how to prevent them

- Public storage buckets: data meant to be private, set to public. Prevent with account-wide public-access blocks and continuous scanning.
- Over-permissive IAM: wildcard permissions and standing admin. Prevent with least privilege and regular access reviews.
- Exposed management ports: SSH/RDP open to 0.0.0.0/0. Prevent with default-deny and bastion/just-in-time access.
- Disabled logging: you cannot detect what you do not record. Prevent by making audit logging mandatory and centralised.
- Secrets in code: keys committed to repositories. Prevent with secret scanning and a secrets manager (Chapter 9).

INSIGHT: Why misconfiguration dominates

Cloud rewards speed, and speed produces mistakes. A developer under deadline clicks 'make public' to unblock a demo and forgets. The fix is not blaming people — it is automation: tools that detect and prevent misconfiguration continuously, so a human slip is caught in minutes, not months. This is what CSPM (Chapter 16) exists to do.

9. Secrets Management

A 'secret' is any credential that grants access: API keys, database passwords, tokens, private keys. Mishandled secrets are a top breach cause because they are the literal keys to the kingdom, and they leak easily — into code, logs, config files, and chat messages.

The lifecycle of a secret

- Never hardcode: secrets in source code end up in version history forever, even if deleted later.
- Store in a secrets manager: a dedicated, encrypted, access-controlled, audited store (e.g. AWS Secrets Manager, Azure Key Vault, HashiCorp Vault).
- Inject at runtime: the application fetches the secret when it runs, rather than shipping with it.
- Rotate regularly: change secrets on a schedule so a leaked one has a short useful life.
- Prefer no secret at all: use workload identity / roles so a service authenticates by what it is, not by a shared password it must store.

```
# BAD: secret baked into code, leaks into git history forever
db_password = "P@ssw0rd_123"

# BETTER: fetch from a secrets manager at runtime
db_password = secrets_manager.get("prod/db/password")

# BEST: no stored secret at all — the workload assumes a role
# and authenticates by its own cloud identity
```

PRACTICE: Scan for leaks continuously

Assume secrets will leak, and detect them fast. Run automated secret scanning on every commit and in your repositories. When a secret leaks, the clock is running — attackers scan public code for keys within minutes. Rotate immediately; do not just delete the commit.

10. Logging, Monitoring & Detection

You cannot defend what you cannot see. Logging records what happened; monitoring watches for trouble; detection turns signals into alerts. In the cloud, the good news is that every action is an API call, so — if you enable it — you can log essentially everything.

The layers of visibility

Layer	What it captures	Vendor example
Control-plane audit log	Every API call: who did what, when	CloudTrail, Azure Activity Log, Cloud Audit Logs
Data-access log	Who read/wrote specific data	S3 access logs, storage logs
Network flow log	What talked to what	VPC Flow Logs, NSG flow logs
Threat detection	Analysed signals & findings	GuardDuty, Defender for Cloud, SCC

Turning logs into defence

- Enable audit logging everywhere, centrally: aggregate logs into one account/workspace the attacker cannot easily tamper with.
- Protect the logs: make them immutable and separate from the workloads they record — attackers delete logs to hide.
- Alert on the signals that matter: root/admin usage, IAM policy changes, disabled logging, access from new regions, mass downloads, failed-auth spikes.
- Practise the response: an alert no one acts on is not detection. Route alerts to people or automation that will respond.

DETECTION TIP: The attacker's first move

Sophisticated attackers try to disable or delete logging early, to blind you. So a change that turns off audit logging should itself be one of your loudest alerts — treat 'logging disabled' as an incident until proven otherwise.

PART IV

Defender: Scale & Automation

"Modern cloud is containers, functions, and pipelines. Securing them by hand does not scale — so defence becomes code."

11. Container & Kubernetes Security

Containers package an app with its dependencies so it runs the same everywhere. Kubernetes (K8s) orchestrates thousands of them. Both add power and a new, layered attack surface. The mental model is the 4 C's: security flows from Code to Container to Cluster to Cloud, and a weakness at any layer undermines the ones above it.

The 4 C's: Code → Container → Cluster → Cloud

Code	SAST, secrets scanning, deps
Container image	Minimal base, scan CVEs, sign
Registry	Access control, trusted images
Orchestrator (K8s)	RBAC, network policy, PSA
Host / node	Hardened OS, runtime security

The layers of container security. Each layer has its own controls; a vulnerable base image or an over-permissive cluster role can undo careful work elsewhere.

What goes wrong, and the controls

Layer	Common weakness	Control
Image	Bloated base image full of CVEs	Minimal base, scan images, sign & verify
Registry	Anyone can pull/push images	Access control, trusted registries only
Cluster (K8s)	Over-broad RBAC, flat pod network	Least-privilege RBAC, network policies
Runtime	Containers run as root, no limits	Non-root, drop capabilities, resource limits
Secrets	Secrets in env vars / images	Mounted from a secrets store, not baked in

CORE IDEA: The container-specific trap

A container is not a strong security boundary by itself — it shares the host kernel. A container running as root that escapes can become root on the node, then the cluster. So 'run as non-root, drop privileges, limit what it can do' is not optional hardening; it is the difference between a contained bug and a cluster takeover.

12. Serverless Security

Serverless (functions-as-a-service, e.g. AWS Lambda) lets you run code without managing servers at all. The provider secures more of the stack — but the shared

responsibility line just moves; it does not vanish. Your code, its permissions, and its inputs are entirely yours.

What changes and what stays

- No servers to patch — the provider handles the runtime and OS. That is a real win.
- Function permissions are the new battleground: each function has an identity/role. Over-privilege it and a code bug becomes a breach — the SSRF chain applies here too.
- Event-driven inputs: functions are triggered by events (an upload, a message, an HTTP request). Every trigger is untrusted input; validate it.
- More functions, more roles, more surface: hundreds of small functions each need least-privilege scoping — the volume is the challenge.

MISCONCEPTION: The serverless myth

“Serverless is secure because there is no server.” Wrong lens. You removed patching, not responsibility. The most common serverless breach is an over-privileged function role plus unvalidated input — exactly the same failure as everywhere else in this course, in a new shape.

13. DevSecOps & Infrastructure as Code

Cloud infrastructure is created by code — Infrastructure as Code (IaC), e.g. Terraform. That is a gift to security: if infrastructure is defined in text, you can review and scan it before it exists. This is ‘shift left’ — catching problems early, when they are cheap to fix.

Security as code, across the pipeline

Stage	Security activity	Example tooling type
Write code	Scan source for bugs & secrets	SAST, secret scanners
Define infra (IaC)	Scan templates for misconfig before deploy	IaC scanners (policy-as-code)
Build	Scan dependencies & images for CVEs	SCA, image scanners
Deploy	Enforce policy gates; block non-compliant	Policy-as-code, admission control
Run	Continuously detect drift & misconfig	CSPM / CNAPP

Why this is the highest-leverage shift

- Cheap to fix early: a misconfiguration caught in a code review costs minutes; the same one in production can cost a breach.
- Consistent and repeatable: policy-as-code enforces the same rules every time, with no human forgetting.
- Auditable: every infrastructure change is a reviewed, version-controlled commit — you can prove what changed and why.

CAREER: The cultural shift

DevSecOps is less a toolset than a principle: security is everyone's job, built into the pipeline, not a gate at the end. The security team's role shifts from 'saying no' late to 'building guardrails' early, so developers move fast and safely.

PART V

Hero: Govern & Operate

"Knowing controls makes you useful. Running a program that proves, governs, and improves them makes you valuable."

14. Compliance & Frameworks

As you grow, you must speak the language of governance — because that is how cloud security gets funded and enforced. You do not need to memorise every clause, but you must know the landmarks and what each gives you.

Framework	What it gives you
CIS Benchmarks	Concrete, per-service hardening checklists (the practical starting point)
CSA Cloud Controls Matrix	A cloud-specific control framework mapped to other standards
NIST CSF / 800-53	A broad risk-and-control language widely used in the US
ISO/IEC 27017	Cloud-specific extension of the ISO 27001 security standard
SOC 2	An attestation of a service provider's controls (often required by customers)
Well-Architected (vendor)	Provider 'security pillar' guidance for their own platform

Compliance is a floor, not a ceiling

A critical distinction: compliance is not security. Passing an audit means you met a minimum bar on the day of the audit. Attackers do not care about your certificate. Use frameworks to structure and prove your work — but treat them as the floor of what you do, never the goal.

TIP: How to actually use a framework

Start with CIS Benchmarks for your provider — they are concrete and immediately actionable. Automate checking against them (most CSPM tools do this out of the box). Then map your controls to a broader framework (NIST, CSA) so you can answer an auditor with evidence and trends, not promises.

15. Cloud Incident Response

Assume breach. A mature program is not one that never has incidents — it is one that detects and contains them fast. Cloud incident response follows the classic phases, but the cloud changes the mechanics: everything is API-driven, which cuts both ways — attackers move fast, but so can your automated response.

The phases, cloud-flavoured

Phase	Cloud specifics
Prepare	Log everything centrally; pre-write playbooks; define who can act
Detect & analyse	Use threat-detection findings and audit logs to scope the incident
Contain	Revoke credentials, isolate resources, tighten security groups — via API, in seconds

Eradicate	Remove attacker access, rotate all potentially exposed secrets, fix root cause
Recover	Rebuild from known-good IaC; restore data; verify integrity
Learn	Post-incident review; turn the incident into new detections and guardrails

SYSTEMS VIEW: The cloud advantage in IR

Because everything is code and API, cloud incident response can be automated: a detection can auto-revoke a credential or quarantine a resource in seconds, faster than any human. Building these automated responses is a hallmark of an advanced program — and a strong interview talking point.

The single most important preparation

Revoke-and-rebuild capability. If you can (a) instantly revoke any credential and (b) rebuild any environment from version-controlled IaC, you can respond to almost any incident with confidence. Both come from the disciplines earlier in this course — least privilege and infrastructure as code — which is why the ladder is ordered the way it is.

16. Building a Cloud Security Program

The final rung: turning individual controls into a running program that scales across teams and improves over time. This is what separates someone who knows cloud security from someone who runs it.

The tooling landscape (know the acronyms)

Acronym	Full name	What it does
CSPM	Cloud Security Posture Management	Continuously finds misconfigurations across accounts
CWPP	Cloud Workload Protection Platform	Protects the workloads themselves (VMs, containers)
CIEM	Cloud Infrastructure Entitlement Mgmt	Finds and right-sizes excess IAM permissions
CNAPP	Cloud-Native Application Protection	Combines the above into one platform

The pillars of a real program

- **Visibility:** know every account, resource, and identity you have. You cannot secure what you do not know exists.
- **Guardrails, not gates:** prevent bad configurations automatically (policy-as-code) rather than reviewing everything manually.
- **Least privilege at scale:** continuously reduce excess permissions (CIEM) — permissions sprawl over time and must be pruned.
- **Continuous detection & response:** assume breach; monitor everything; automate containment.

- Measure and report: track posture over time and map it to a framework so leadership sees progress, not just activity.

THE SUMMIT: From technician to leader

The difference between knowing attacks and running a program is governance: ownership, automation, measurement, and continuous improvement. A hero does not personally fix every misconfiguration — they build the system that prevents, detects, and fixes them at scale. That is the real meaning of the top rung.

PART VI

Interview Preparation

“Understanding is proven under pressure. These questions turn what you have learned into what you can defend out loud.”

17. Interview Q&A Bank

Sixty-plus graded questions, organised by difficulty. Each answer is a model answer — study the reasoning, then rephrase it in your own words. A memorised answer fails the first follow-up; an understood one survives ten. Tags: [F] Foundational, [P] Practitioner, [A] Advanced.

Foundations

[F] Q1. Explain the Shared Responsibility Model in one breath.

The provider secures the cloud itself (physical facilities, hardware, hypervisor); the customer secures what they put in the cloud (data, access, configuration). The exact line shifts by service model — more is the provider's job as you go IaaS → PaaS → SaaS — but data and access are always the customer's responsibility.

[F] Q2. Why is 'identity is the new perimeter' true in the cloud?

Because cloud resources are reachable over the internet via API by default, there is no network wall to get through. An attacker mainly needs valid credentials or a misconfiguration. So who you are (identity and its permissions) becomes the primary control, replacing where you are (network location).

[F] Q3. What is the difference between IaaS, PaaS, and SaaS from a security view?

They differ in how much you must secure. IaaS: you secure the OS, apps, data, and access. PaaS: the provider manages the platform, you secure your code, data, and access. SaaS: you mainly secure your data and who can access it. The customer's share shrinks but never reaches zero.

[F] Q4. What is least privilege and why does it matter so much in the cloud?

Granting an identity only the exact permissions its task requires — nothing more. It matters because cloud credentials are powerful and reachable via API; a stolen over-privileged credential can rewrite the whole account, while a least-privileged one caps the damage to almost nothing.

[F] Q5. Name the most common cause of cloud data breaches.

Customer misconfiguration — for example public storage buckets, over-permissive IAM, or exposed management ports — not provider failure. It dominates because the cloud exposes thousands of settings and rewards speed, so human mistakes are frequent and instantly internet-facing.

[F] Q6. Encryption at rest vs. in transit?

At rest protects stored data on disk, so a stolen disk or storage object is unreadable. In transit protects data moving over the network (TLS), so it cannot be read if intercepted. You need both; neither protects against a valid credential that is simply allowed to read the data.

[F] Q7. Why is encryption not enough on its own?

Because encryption protects against theft of the raw bytes (a stolen disk), not against an authorised identity reading the data through the normal path. Most breaches are

access-control failures, not broken crypto. Encryption and IAM are partners.

[F] Q8. What is MFA and where is it non-negotiable?

Multi-factor authentication requires a second proof of identity beyond a password. It is non-negotiable on all privileged and administrative accounts, and strongly recommended everywhere, because a password alone is one phishing or leak away from full compromise.

[F] Q9. What is a security group?

A stateful virtual firewall attached to a resource that controls inbound and outbound traffic by port, protocol, and source. Best practice is default-deny: allow only required traffic, and never expose management ports like SSH or RDP to the entire internet.

[F] Q10. Compliance vs. security — are they the same?

No. Compliance means meeting a defined minimum standard, verified at a point in time; security is the ongoing practice of protecting systems against real attackers. Compliance is a useful floor and a way to prove work, but passing an audit does not mean you are secure.

[F] Q11. What is the CIA triad, cloud-flavoured?

Confidentiality (keep data secret), Integrity (keep data and config correct), Availability (keep systems and budget usable). The cloud twist: availability includes financial availability — an attacker can inflict harm by running up a massive bill.

[F] Q12. What does ‘default deny’ mean?

Block all traffic and access by default, then explicitly allow only what is needed. It is safer than ‘allow all, then block bad’ because anything you forget stays closed rather than exposed.

Practitioner

[P] Q13. Walk through the SSRF-to-credential-theft attack chain.

An attacker exploits a server-side request forgery bug to make the app fetch the cloud metadata endpoint, which returns the workload’s temporary IAM credentials. If the workload’s role is over-privileged, the attacker uses those credentials to access and exfiltrate data — for example from storage. Defences: validate inputs, harden the metadata service, apply least privilege to the role, segment the network, and monitor data access.

[P] Q14. Your storage bucket was found public. Walk through your response.

Contain: remove public access immediately and enable account-wide public-access blocking. Assess: use access logs to determine what was exposed and whether it was accessed. Notify per policy and any legal obligations. Fix root cause: find why it was public (a script, a manual click), add a preventive guardrail, and add continuous scanning so it cannot recur silently.

[P] Q15. How do you enforce least privilege at scale across hundreds of identities?

Start every identity from zero and add only needed permissions. Use CIEM tooling to detect and remove excess entitlements continuously, review access regularly, prefer temporary roles over long-lived keys, and use policy-as-code to prevent wildcard permissions from being deployed in the first place.

[P] Q16. Why prefer IAM roles over long-lived access keys?

Roles issue temporary, automatically rotating credentials scoped to a task, so a leak has a short useful life and limited scope. Long-lived keys sit in files and configs indefinitely, are easy to leak, and rarely get rotated — making them a favourite target.

[P] Q17. How would you secure secrets in a CI/CD pipeline?

Never hardcode them. Store them in a secrets manager, inject them at runtime with least-privilege access, rotate them regularly, and run secret scanning on every commit to catch leaks. Where possible, use workload identity so the pipeline authenticates by what it is rather than storing a shared secret.

[P] Q18. What logs would you enable first in a new cloud account, and why?

The control-plane audit log (who made which API call, when) first — it is the backbone of both detection and incident response. Then data-access logs and network flow logs. Centralise them into a separate, tamper-resistant location, because attackers try to delete logs to hide.

[P] Q19. An alert fires that audit logging was disabled. What do you think and do?

Treat it as a probable incident until proven otherwise, because disabling logging is a classic attacker move to gain blindness. Investigate who made the change and why, re-enable logging, and check for other suspicious activity around the same time and identity.

[P] Q20. How do you limit lateral movement in a cloud network?

Segment: separate environments and tiers into distinct networks/subnets with default-deny rules, so a compromise in one cannot freely reach others. Keep databases and internal services in private subnets, use private endpoints, and apply Zero Trust — verify identity on every request rather than trusting the internal network.

[P] Q21. What are the 4 C's of container security?

Code, Container, Cluster, Cloud. Security flows through these layers: secure the application code and dependencies, build minimal scanned and signed images, harden the orchestrator with least-privilege RBAC and network policies, and secure the underlying cloud. A weakness at any layer undermines the ones above.

[P] Q22. Why is a container not a strong security boundary?

Because containers share the host's kernel. A container running as root that exploits a kernel or misconfiguration flaw can escape to the host node and then the cluster. That is why running as non-root, dropping capabilities, and setting resource limits are essential, not optional.

[P] Q23. What is the most common serverless security failure?

An over-privileged function role combined with unvalidated event input. Serverless removes server patching but not responsibility: the function's identity, its permissions, and its inputs are still yours, and the same least-privilege-plus-input-validation discipline applies.

[P] Q24. What does 'shift left' mean and why is it high-leverage?

Moving security earlier in the development lifecycle — scanning code and infrastructure-as-code before deployment. It is high-leverage because a misconfiguration caught in code review costs minutes, while the same one in production can cost a breach;

and policy-as-code enforces rules consistently.

[P] Q25. How does Infrastructure as Code help security?

It turns infrastructure into reviewable, version-controlled text you can scan before anything is created. That enables catching misconfigurations pre-deployment, enforcing policy-as-code, auditing every change, and rebuilding cleanly from a known-good definition during incident recovery.

[P] Q26. Name three signals worth alerting on and why.

Root/admin account usage (should be rare; often attacker activity), IAM policy changes (privilege escalation), and logging being disabled (attacker gaining blindness). Each is a strong early indicator of compromise with low false-positive cost.

Advanced

[A] Q27. 'It's in the cloud, so the provider secures it.' Respond.

Partly true, dangerously incomplete. The provider secures their layer — hardware, facilities, hypervisor — under the Shared Responsibility Model. But the customer always owns data, access, and configuration, and misconfiguration by customers is the leading breach cause. Security in the cloud is a shared job, and the customer's half is where most breaches occur.

[A] Q28. Design a secure baseline for a brand-new cloud account.

Enable centralised, immutable audit logging from day one. Enforce MFA on all privileged identities and remove standing admin. Apply least privilege and block wildcard policies via policy-as-code. Block public storage access account-wide. Segment networks with default-deny. Enable threat detection and route findings to real responders. Enforce encryption at rest and in transit. Manage all infrastructure as reviewed IaC and scan it pre-deploy.

[A] Q29. Where is the highest-leverage single control in the SSRF breach chain, and why?

Least privilege on the workload role. Every other defence breaks one link, but least privilege caps the impact of the whole chain: even if SSRF succeeds and credentials are stolen, a tightly scoped role means the attacker can reach almost nothing. It converts a potential full breach into a contained bug.

[A] Q30. How would you measure the security posture of a large cloud estate?

Continuously scan all accounts with CSPM against a concrete benchmark (e.g. CIS) and track the misconfiguration count and its trend over time. Add coverage metrics: percentage of identities at least privilege (via CIEM), percentage of workloads with logging and encryption enabled, and mean time to detect and contain incidents. Report trends mapped to a framework, not one-off snapshots.

[A] Q31. Explain defence in depth using a real cloud breach pattern.

Take the SSRF-to-exfiltration chain: input validation, metadata hardening, least privilege, network segmentation, and monitoring are five independent layers. No single one stops every variant, but any one breaks the specific chain. Defence in depth assumes each control can fail and ensures another stands behind it, so a single failure is never catastrophic.

[A] Q32. What new risks do multi-cloud environments introduce?

Inconsistent controls and terminology across providers, a larger and more fragmented attack surface, identity sprawl across separate IAM systems, and gaps where teams know one cloud well and another poorly. Mitigate with a unified control framework mapped to each provider, centralised visibility (CNAPP), and consistent policy-as-code rather than per-cloud manual configuration.

[A] Q33. How do you build automated incident response in the cloud?

Wire detections to actions: a threat-detection finding or a specific log event triggers automation that, for example, revokes a credential, quarantines a resource, or tightens a security group in seconds. Pre-write and test these playbooks, scope the automation's own permissions carefully, and keep a human in the loop for high-impact or ambiguous cases.

[A] Q34. Argue whether compliance certification makes a company secure.

It does not, though it helps. Certification proves a minimum set of controls existed at audit time; it is a floor and an evidence tool, not proof against real attackers, who ignore certificates. A certified company with an over-privileged role and a public bucket is still breached. Use compliance to structure and demonstrate a program, then exceed it.

[A] Q35. A developer needs broad access 'to move fast.' How do you respond?

Reframe from access to outcome: identify the specific actions the task needs and grant exactly those, just-in-time, rather than standing broad access. Explain the risk in their terms — a bug in their internet-facing code plus a broad role equals a full-account breach. Offer guardrails (policy-as-code, temporary elevation) that let them move fast safely, so security enables rather than blocks.

[A] Q36. What is your one-sentence philosophy of cloud security?

Assume the perimeter is gone, so enforce security through identity (least privilege), configuration (automated guardrails), and visibility (log and detect everything) — and assume breach, so that when one control fails, another contains it.

Rapid-fire & scenario

[F] Q37. What is a VPC?

A virtual private cloud: your own logically isolated network within the provider, where you define subnets, routing, and firewall rules to control what can reach what.

[F] Q38. What is CSPM?

Cloud Security Posture Management: tooling that continuously scans your cloud accounts for misconfigurations and compliance violations, so mistakes are caught in minutes rather than months.

[F] Q39. What is a bastion / jump host?

A hardened, tightly controlled entry point used to reach private resources, so management access (SSH/RDP) is not exposed directly to the internet.

[F] Q40. What is the principle behind Zero Trust?

Never trust, always verify: authenticate and authorise every request regardless of network location, rather than assuming anything inside the network is safe.

[F] Q41. What is data classification and why do it?

Labelling data by sensitivity (public, internal, confidential, regulated) so that controls match the risk — you protect regulated PII far more strictly than public marketing content.

[F] Q42. What is CIEM?

Cloud Infrastructure Entitlement Management: tooling that discovers and right-sizes IAM permissions, finding and removing the excess entitlements that accumulate over time.

[P] Q43. How do you protect the cloud metadata endpoint?

Enforce the hardened, session-based version of the metadata service (which defeats simple SSRF), apply least privilege to the workload role so stolen credentials are near-useless, and block the app from making requests to internal IP ranges.

[P] Q44. How would you detect crypto-mining in your account?

Watch for anomalous compute spikes, unexpected instance launches (especially large GPU types), traffic to known mining pools, and sudden cost increases. Threat-detection services flag several of these; a billing anomaly alert is a simple, effective backstop.

[P] Q45. What is policy-as-code and where does it run?

Expressing security and compliance rules as code that is automatically evaluated — for example scanning IaC before deployment and blocking non-compliant changes at an admission gate. It enforces rules consistently without relying on humans to remember them.

[P] Q46. How do you handle a leaked long-lived access key?

Immediately deactivate and delete the key, investigate what it accessed via audit logs, rotate anything it could have exposed, and determine how it leaked (e.g. committed to code). Then reduce reliance on long-lived keys by moving to temporary role-based credentials.

[A] Q47. Design monitoring for a regulated workload holding PII.

Enable control-plane, data-access, and network flow logging, centralised and immutable. Alert on any access to the PII data store from unusual identities, locations, or at unusual volume; on IAM changes touching that data; and on logging or encryption being disabled. Enforce encryption with customer-managed keys and log every key use, so you can prove who could decrypt what.

[A] Q48. When would you accept a security risk rather than fix it?

When the cost or friction of the control clearly exceeds the expected loss (likelihood times impact), and a decision-owner formally accepts it. Risk acceptance is legitimate but must be explicit, time-bounded, documented, and revisited — not a silent default. The job is informed risk management, not zero risk.

That is 48 graded questions plus rapid-fire. Extend the bank yourself: for every control in this course, write a 'how would you test it works' question, and for every attack, a 'how would you detect and contain it' question. Teaching each answer aloud is the fastest way to expose gaps.

METHOD: How to rehearse this bank


Cover the answer, read the question, and speak for 30–60 seconds before revealing the model answer. Score yourself on three things: correct core idea, handling of the obvious follow-up, and a concrete example. Any question missing all three points to a chapter to reread. Re-test weak questions after two days — spacing is what makes recall durable.

18. Course vs. Q&A: Which Matters More?

You asked at the very start which is more important — the concept course or the interview questions. Here is the reasoned answer, because the question deserves one.

Reframing the question

“Which is more important” assumes they compete. They do not; they do different jobs. The course builds a mental model. The Q&A tests and signals that model. Asking which matters more is like asking whether training or the exam matters more — the exam is worthless without the training, and the training is unverified without the exam.



The recommended split of study time. Understanding is the larger investment; rehearsal converts it into interview performance.

The argument, cause to effect

- Cause: cloud interviews probe with ‘why?’ and ‘what if?’ Effect: memorised answers without a model collapse; understood answers branch into correct follow-ups.
- Cause: the job is novel-problem-solving (new services appear constantly), not recall. Effect: employers screen for reasoning, so the course’s mental models are what get you hired and keep you employed.
- Cause: a question bank is finite; real misconfigurations are infinite. Effect: only a transferable model generalises to problems no bank contains.

But Q&A is not optional

Rehearsal matters because understanding does not automatically become fluent speech under pressure. The Q&A bank trains retrieval and articulation — turning what you know into what you can say crisply. Skipping it means understanding the material but freezing when asked.

ANSWER: The verdict

Spend about 70% of your effort understanding (Parts I–V) and 30% rehearsing Q&A (Part VI). If forced to pick one, the course wins — you can reason your way to an answer you never rehearsed, but you cannot rehearse your way to reasoning you never built. Use the Q&A as a diagnostic: every question you cannot answer points to a chapter to reread.

Appendix A — 30-Day Study Plan

A concrete schedule enforcing the 70/30 split. Adjust the pace to your background, but keep the order — the ladder is deliberate.

Days	Focus	Outcome you should reach
1–3	Part I (Foundations)	Draw the Shared Responsibility Model from memory.
4–9	Part II (IAM, network, data)	Explain least privilege and encryption states clearly.
10–16	Part III (Attacks & defence)	Trace the SSRF chain and name a defence per link.
17–22	Part IV (Containers, serverless, DevSecOps)	Explain the 4 C’s and shift-left.
23–26	Part V (Govern & operate)	Design a secure baseline and an IR playbook.
27–30	Part VI (Q&A rehearsal)	Answer every question aloud; reread weak chapters.

LEARNING SCIENCE: Active recall beats rereading

After each part, close the notes and write what you remember before checking. The effort of retrieval builds durable memory; passive rereading feels productive but is not. Pair it with hands-on practice in a free-tier cloud account — reading about a public bucket is nothing like fixing one.

Appendix B — Glossary

Term	Definition
Bastion host	A hardened entry point for reaching private resources without exposing them directly.
CIEM	Tooling to find and right-size excess IAM permissions.
CNAPP	A platform combining posture, workload, and entitlement protection.
CSPM	Continuous scanning for cloud misconfigurations and compliance drift.
Default deny	Block everything, then explicitly allow only what is needed.

Encryption at rest	Encrypting stored data so a stolen disk or object is unreadable.
Encryption in transit	Encrypting data moving over the network (TLS).
IaaS / PaaS / SaaS	Service models defining how much of the stack you rent — and secure.
IAM	Identity and Access Management: who can do what to which resource.
IaC	Infrastructure as Code: defining cloud infrastructure in reviewable text.
KMS	Key Management Service: stores, controls, and audits encryption keys.
Least privilege	Granting only the exact permissions a task requires.
Metadata endpoint	A cloud-internal service that hands a workload its temporary credentials.
MFA	A second proof of identity beyond a password.
Misconfiguration	An insecure setting; the leading cause of cloud breaches.
Security group	A stateful virtual firewall around a resource.
Shared Responsibility	The split of security duties between provider and customer.
SSRF	Server-Side Request Forgery: tricking a server into making attacker-chosen requests.
VPC / VNet	Your isolated virtual network in the cloud.
Zero Trust	Never trust, always verify — authorise every request regardless of location.

End of course. You started at zero. If you can now draw the Shared Responsibility Model, trace the SSRF breach chain, design a secure baseline, and defend the Part VI answers under follow-up questioning, you have reached hero. Keep practising in a real account — cloud security is a skill, not a syllabus.

nexroa · cybersecurity education